

# OS HW2 report

Student ID : 314551123 / Name : 蘇秭郁

---

- Answer the Q&As

1. Describe how you implemented the program in detail.
  - (1) First, **getopt()** is utilized to parse the input arguments, and **strup()** is employed to back up the parameter strings; this prevents **strtok\_r()** from directly modifying the original content. Subsequently, **sched\_setaffinity** is used to bind the process to a single CPU, which is essential for observing preemption and Round-Robin (RR) scheduling behaviors.
  - (2) Next, the corresponding threads are created based on the input parameters:
    - i. When configuring thread\_attr, **pthread\_attr\_setinheritsched(&attr, PTHREAD\_EXPLICIT\_SCHED)** is invoked to ensure that the created threads do not inherit attributes from the main thread, but instead use the specifically defined attributes.
    - ii. For the priority of SCED\_FIFO threads, the priority levels are configured via the **sched\_param** structure.
    - iii. To retrieve policies and priorities, **strtok\_r()** is used instead of the standard strtok() to avoid thread-local storage conflicts and potential segmentation faults.
  - (3) Furthermore, to ensure that all threads begin competing for the CPU at the same time, a **pthread\_barrier\_t** is implemented. After each thread is created, it calls **pthread\_barrier\_wait**, remaining blocked until all threads have been successfully initialized, at which point they begin execution simultaneously.
  - (4) Regarding time measurement, since the busy-wait period must exclude the time the thread is preempted, **clock\_gettime()** is used in conjunction with **CLOCK\_THREAD\_CPUTIME\_ID**. This specific timer ensures that counting only occurs when the thread is actively utilizing the CPU.

Within the program, the CPU occupancy time is calculated through a busy-wait loop.

- (5) Finally, **pthread\_join()** is called within main() to wait for all worker threads to complete their execution. Once finished, the barrier is destroyed, and the allocated memory is released.
  
2. Describe the results of `./sched_demo -n 3 -t 1.0 -s NORMAL,FIFO,FIFO -p -1,10,30` and what causes that.

```
~ # ./sched_demo -n 3 -t 1.0 -s NORMAL,FIFO,FIFO -p -1,10,30
Thread 2 is running
Thread 2 is running
Thread 2 is running
Thread 1 is running
Thread 1 is running
Thread 1 is running
Thread 0 is running
Thread 0 is running
Thread 0 is running
```

Since Thread 1 and Thread 2 are assigned the FIFO scheduling policy while Thread 0 is set to NORMAL, Thread 1 and Thread 2 will be executed with higher priority. Furthermore, because Thread 2 has a higher priority level than Thread 1, Thread 2 will execute before Thread 1. Consequently, the final execution sequence will be Thread 2 → Thread 1 → Thread 0.

3. Describe the results of `./sched_demo -n 4 -t 0.5 -s NORMAL,FIFO,NORMAL,FIFO -p -1,10,-1,30`, and what causes that.

```
~ # ./sched_demo -n 4 -t 0.5 -s NORMAL,FIFO,NORMAL,FIFO -p -1,10,-1,30
Thread 3 is running
Thread 3 is running
Thread 3 is running
Thread 1 is running
Thread 1 is running
Thread 1 is running
Thread 2 is running
Thread 0 is running
Thread 2 is running
Thread 0 is running
Thread 2 is running
Thread 0 is running
```

Since Thread 1 and Thread 3 are assigned the FIFO scheduling policy, they will execute before Thread 0 and Thread 2. Within the FIFO group, Thread 3 has a higher priority than Thread 1, so Thread 3 will execute first. The remaining Thread 0 and Thread 2 engage in time-sharing execution, where the Completely Fair Scheduler (CFS) manages their resource allocation in an

interleaved fashion.

4. Describe how did you implement n-second-busy-waiting?

The n-second busy waiting is implemented using `clock_gettime()` with the `CLOCK_THREAD_CPUTIME_ID` clock source. This timer measures the actual CPU time consumed by the thread, effectively excluding any duration where the thread is preempted. By polling the timer within a while loop without yielding the CPU, the program ensures the thread remains busy for exactly the specified duration.

5. What does the `kernel.sched_rt_runtime_us` effect? If this setting is changed (eg. 500000, 950000, 1000000), what will happen?

The `kernel.sched_rt_runtime_us` parameter is a safety mechanism in Linux designed to prevent real-time threads from completely starving the CPU. It specifies the amount of time real-time (RT) threads are allowed to run within a specific period (defined by `sched_rt_period_us`, which defaults to 1 second).

- Effects of different settings:

500,000 (0.5s): RT threads are capped at 50% CPU utilization per second. The remaining 50% is strictly reserved for non-RT tasks (e.g., normal processes, shell commands).

950,000 (0.95s, Default): This is the system default. RT threads can consume up to 95% of the CPU. The 5% margin ensures that a system administrator can still access a shell to kill a runaway or bugged RT process.

1,000,000 (1s): RT threads can occupy 100% of the CPU. This is dangerous because if a `SCHED_FIFO` process enters an infinite loop, the system will become completely unresponsive, preventing any keyboard or mouse input.